

SmartShield — Sprint 0 Research & Contribution Report

Author: Aditya Pathak

Role: Research, Dataset & ML Preparation (MEMBER 3)


Sprint: Sprint 0 (16 August 2026 – 23 August 2026)

Deliverables Covered: [Vulnerability_Research.md](#), [Dataset_Analysis.md](#), [ML_Baseline.md](#), [API_Integration.md](#)

Status: Completed & Revised for Sprint 1 Readiness

1. Executive Summary & Role Scope

During Sprint 0, my core objective was to eliminate architectural uncertainty for SmartShield by establishing the empirical and technical foundation for the entire project. In the project dependency chain, vulnerability research and dataset planning serve as the primary upstream drivers:



By defining the MVP vulnerability scope, designing a leak-proof dataset evaluation pipeline, and formalizing the ML baseline and API contracts, my deliverables provide direct specifications for Intermediate Representation (IR) design, graph construction, and deterministic detector implementation.

2. Updated Deliverables Summary

A. Vulnerability Scope & Prioritization ([Vulnerability_Research.md](#))

- **Structured Scoring Model:** Implemented an objective 4-metric evaluation framework (Static-Analysis Suitability, Dataset Availability, False Positive Risk, Demonstration Value) to evaluate 8 candidate vulnerability classes.
- **MVP Scope Lock:** Confirmed **Reentrancy** and **tx.origin Misuse** as the two primary MVP targets, providing high static-analysis viability and sufficient dataset support.
- **Compiler-Aware Scoping:** Formally rejected Integer Over/Underflow from the MVP scope to prevent downstream architectural rework, accounting for native Solidity **0.8.0+** arithmetic checks.
- **False-Positive Mitigation:** Formulated precise detection logic requirements for downstream detectors, distinguishing between benign AST occurrences and exploitable control-flow patterns.

B. Dataset Investigation & De-Duplication ([Dataset_Analysis.md](#))

- **Dataset Auditing:** Analyzed the structural profiles of **SolidiFI** (synthetic bug injection) and **SmartBugs Curated** (historical ground truth).
- **Data Leakage Prevention:** Established explicit protocols for AST-level de-duplication and contract-level splitting to prevent training contamination across isomorphic smart contracts.

C. Machine Learning Baseline Architecture ([ML_Baseline.md](#))

- **IR-Coupled Features:** Designed feature vectors derived directly from SmartShield's Intermediate Representation and Control Flow Graphs rather than raw source code.
- **Two-Experiment Evaluation Strategy:** Formalized a dual-track testing protocol separating within-dataset generalization (SolidiFI 80/20 split) from cross-dataset zero-shot generalization (train on SolidiFI, test on SmartBugs Curated).
- **Strict ML Retention Gate:** Set a quantitative success criterion requiring the ML model to demonstrate measurable value over deterministic rule-based detectors on real-world test sets before permanent adoption.

D. System Boundaries & Decoupling ([API_Integration.md](#))

- **Stateless Execution:** Defined input/output requirements ensuring the C++ analysis engine functions as an independent, stateless analyzer.
- **Adapter Pattern Isolation:** Decoupled internal C++ data structures from external presentation layers using a versioned JSON schema and an explicit serialization adapter layer.

3. Technical Research & Methodological Foundations

3.1 Vulnerability Mechanics & Static Analysis Requirements

1. Reentrancy

- **Exploitation Pattern:** Occurs when a contract performs an external call (e.g., `.call{value: ...}("")`) before updating its internal accounting state (violating the Checks-Effects-Interactions pattern). A malicious contract intercepts control flow via its

`fallback()` or `receive()` function and invokes the calling function repeatedly before the initial execution completes.

- **Required Static Analysis Context:** The static analysis engine cannot rely on pattern matching alone. It requires:
 - **Call Graph:** To track external calls and identify call target destinations.
 - **Control Flow Graph (CFG):** To determine execution ordering between external call nodes and state-variable write nodes.
 - **Modifier Awareness:** To identify synchronization guards (e.g., OpenZeppelin's `ReentrancyGuard` / `nonReentrant`) and prevent false positives on guarded functions.

2. `tx.origin` Misuse (Occurrence vs. Contextual Misuse)

- **Exploitation Pattern:** `tx.origin` refers to the original external account (EOA) that signed and initiated the transaction call chain, whereas `msg.sender` refers to the immediate caller. In phishing scenarios, an attacker tricks an authorized account into executing a transaction on a malicious contract, which then calls the vulnerable contract; if authorization checks `tx.origin == owner`, the check passes maliciously.
- **Required Static Analysis Context:**
 - *Naive detection (AST search):* Flags any instance of `tx.origin`, producing high false positives on benign uses (such as logging or `require(tx.origin == msg.sender)` anti-contract checks).
 - *SmartShield logic:* Traverses the AST/CFG to verify if `tx.origin` operates within a conditional statement (`require`, `assert`, `if`) that acts as a branch guard for state-changing operations or asset transfers.

3. Arithmetic Semantics & Solidity Compiler Evolution

- **Research Finding:** In Solidity versions prior to `0.8.0`, arithmetic operations (`+`, `-`, `*`) wrap silently on overflow/underflow, requiring explicit safety libraries like `SafeMath`. However, starting in Solidity `0.8.0` (December 2020), arithmetic operations include built-in overflow checks that automatically revert execution with Panic error code `0x11`.
- **Architectural Justification:** Treating "absence of `SafeMath`" as a vulnerability is obsolete for modern codebases. Introducing compiler-version-sensitive parsing in Sprint 0 would overcomplicate the IR; therefore, arithmetic analysis is cleanly excluded from the MVP scope.

3.2 Dataset Selection & Leakage Mitigation

DATASET EVALUATION FRAMEWORK

SolidiFI Benchmark	SmartBugs Curated
• ~9,369 injected vulnerabilities	• ~143 curated contracts
• Systematic bug injection rules	• Real-world historical exploits
• Role: Feature Training / Split	• Role: Zero-Shot Benchmark

Dataset Characteristics

1. **SolidiFI Benchmark:** An academic benchmark utilizing automated bug-injection engines across clean base contracts. It provides exact ground-truth injection coordinates across multiple vulnerability classes, making it ideal for supervised training volume.
2. **SmartBugs Curated Dataset:** A benchmark of 143 historical contracts manually validated and mapped to the DASP taxonomy. It captures organic, developer-written vulnerabilities rather than templated injections.

Mitigating Dataset Leakage in Smart Contract Codebases

Smart contract repositories frequently contain duplicated token templates and cloned contract architectures. Standard random splitting (k -fold or random 80/20) causes data leakage where identical function bodies appear in both training and test partitions.

- **AST-Level Hashing:** Base contract structures are hashed to identify and remove isomorphic clones before partitioning.
- **Strict Contract-Level Partitioning:** Partitioning is performed exclusively at the contract level, ensuring no functions from a training contract are evaluated within the test set.

3.3 Machine Learning Baseline Rationale

Tabular Random Forest vs. Deep Learning / LLMs


In alignment with the project's Sprint 0 constraints (excluding ungrounded neural architectures or LLMs), the ML baseline is built on standard tabular classification ([scikit-learn](#)):

- **Feature Coupling:** Rather than using tokenized source code or NLP embeddings, features are extracted directly from the SmartShield Intermediate Representation (IR) and CFG node attributes (e.g., count of state writes following external calls, branch depth, modifier presence).
- **Interpretability:** Tree-based models allow feature importance ranking (MDI / permutation importance), making it clear which IR attributes contribute to vulnerability classification.
- **Generalization Proof (Cross-Dataset Evaluation):** The model is trained on synthetic SolidiFI data and tested on real-world SmartBugs Curated data. This guarantees that high test metrics reflect actual vulnerability detection capability rather than overfitting to synthetic injection templates.

3.4 API Integration & Engine Isolation

Adapter Pattern Architecture

To ensure the internal C++ core analysis engine is not coupled to external presentation tools or backend interfaces, the communication boundary is strictly isolated:



- **Resilience to Engine Refactoring:** If the internal C++ data representations change during Sprint 1 or Sprint 2, only the internal serialization adapter needs adjustment; external API consumers and reporting pipelines remain unaffected.
- **Stateless Operation:** The analysis engine takes source paths and compiler flags as input, executes analysis in memory, and emits standardized JSON to `stdout` or file output without managing databases, server processes, or authentication.

4. Cross-Module Impact & Sprint 1 Readiness

Module Owner	Downstream Dependency	Unblocked By My Sprint 0 Deliverable
Shermale (IR & Interfaces)	Needs explicit field requirements for IR v1.	Reentrancy & <code>tx.origin</code> scope dictates that IR v1 must preserve external call

		nodes, post-call state modifications, and guard conditions.
Parit (CFG & Call Graphs)	Needs execution path modeling requirements.	Defines the necessity of interprocedural call edges and variable write tracking in CFG nodes.
Aayush (Detectors & Tests)	Needs formal detector specifications and test expectations.	Establishes exact detection logic and false-positive edge cases for rule-based implementations.

This document summarizes all Sprint 0 research, dataset preparations, and interface definitions, establishing complete readiness for implementation in Sprint 1.